

Clasificación con red neuronal

June 28, 2016

```
In [11]: %pylab inline
import numpy as np
import pandas as pd

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.regularizers import l1

import pyprind
```

Populating the interactive namespace from numpy and matplotlib

1 Loading data

```
In [2]: data_train = pd.read_csv('UCI HAR Dataset/train/train.csv')
data_test = pd.read_csv('UCI HAR Dataset/test/test.csv')
```

```
In [3]: X_train = data_train.iloc[:, 1:-6]
Y_train = data_train.iloc[:, -6:]
```

```
X_test = data_test.iloc[:, 1:-6]
Y_test = data_test.iloc[:, -6:]
```

```
In [4]: X_train.shape
```

```
Out[4]: (7352, 561)
```

```
In [5]: np_X_train = X_train.as_matrix()
np_Y_train = Y_train.as_matrix()
```

```
np_X_test = X_test.as_matrix()
np_Y_test = Y_test.as_matrix()
```

2 Creating the model

```
In [7]: def create_model(p=None):
        """
        Creates a model based on parameters 'p'.

        Example:

        p = {'hidden': [64,32],
            'activation_func': 'tanh', #'relu', 'tanh',
```

```

        'l1_reg': 0.01,
        'optimizer_method': 'SGD'} #'SGD', 'adadelata'

    model = create_model(p)

"""

if p==None:
    p = {'hidden': [32],
        'activation_func': 'tanh',
        'l1_reg': 0.01,
        'optimizer_method': 'SGD'}
hidden = p.get('hidden', [32])
activation_func = p.get('activation_func', 'tanh')
l1_reg = p.get('l1_reg', 0.01)
optimizer_method = p.get('optimizer_method', 'SGD')

model = Sequential()
model.add(Dense(hidden[0],
                input_dim=561,
                activation=activation_func,
                W_regularizer=l1(l1_reg)))
hidden = hidden[1:]
for h in hidden:
    model.add(Dense(h, activation=activation_func, W_regularizer=l1(l1_reg)))

model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer=optimizer_method, metrics=['accuracy'])

return model

```

3 Training

3.0.1 Cross-validation

```

In [8]: def cross_validation(X,Y, k=10, nn_p=None, verbose=True):
        """
        Runs a k-folded cross validation over a training set. Returns the classification error obtained.

        Example:

            error = cross_validation(np_X_train, np_Y_train)

        """

    if nn_p==None:
        nn_p = {'hidden': [32],
            'activation_func': 'tanh', #'relu', 'tanh',
            'l1_reg': 0.01,
            'optimizer_method': 'SGD', #'SGD', 'adadelata'
            'epochs': 10} #Not used in this example, but in cross-validation.
    hidden = nn_p.get('hidden', [32])
    activation_func = nn_p.get('activation_func', 'tanh')

```

```

l1_reg = nn_p.get('l1_reg', 0.01)
optimizer_method = nn_p.get('optimizer_method', 'SGD')
epochs = nn_p.get('epochs', 10)

groups = np.random.choice(X.shape[0],X.shape[0], replace=False) % k

classification_error = []
if verbose:
    print("{}-folded cross-validation:".format(k))
    bar = pyprind.ProgBar(k)

for i in xrange(k):
    i_train = groups!=i
    i_test = groups==i

    model = create_model(nn_p)
    training_results = model.fit(X[i_train,:], Y[i_train,:],
                                validation_data=(X[i_test:],Y[i_test:]),
                                nb_epoch=epochs,
                                verbose=0)
    classification_error.append(1 - training_results.history['val_acc'][-1])
    if verbose: bar.update()

classification_error = mean(classification_error)
if verbose: print("{}-folded cross-validation classification error: {}".format(k, classification_error))

return classification_error

```

3.0.2 Grid search

```

In [9]: def grid_search(data,
                        hiddens=None, l1_regs=None, activation_func=None, optimizer_method=None, epochs=None,
                        cv_groups=10):
    """
    Runs a grid search using cross-validation over the hyper-parameters passed. Returns
    classification error for each configuration.

    Example:

        hiddens = [[10], [32]]
        l1_regs = [1, 0]
        grid_search_results = grid_search((np_X_train, np_Y_train),
                                          hiddens=hiddens, l1_regs=l1_regs,
                                          cv_groups=2)

    """
    if hiddens==None: hiddens=[32]
    if l1_regs==None: l1_regs=[0.01]
    if activation_func==None: activation_func=['tanh']
    if optimizer_method==None: optimizer_method=['SGD']
    if epochs==None: epochs=[10]

    grid = [{'hidden':a,
              'l1_reg':b,

```

```

        'activation_func':c,
        'optimizer_method':d,
        'epochs':e}
    for a in hiddens
    for b in l1_regs
    for c in activation_func
    for d in optimizer_method
    for e in epochs]

    bar = pyprind.ProgBar(len(grid))
    for p in grid:
        p['classification_error'] = cross_validation(data[0], data[1], nn_p=p, k=cv_groups, ver
        bar.update()

    return grid

In [12]: # Hacemos una grid search
hiddens = [[64,32], [32,10], [10,5], [64], [32], [10]]
l1_regs = [1, 1e-1, 1e-2, 1e-3, 1e-4, 0]
activation_func = ['tanh']
optimizer_method = ['SGD', 'adadelata']
epochs = [10]

grid_search_results = grid_search((np_X_train, np_Y_train),
                                   hiddens=hiddens,
                                   l1_regs=l1_regs,
                                   activation_func=activation_func,
                                   optimizer_method=optimizer_method,
                                   epochs=epochs,
                                   cv_groups=2)

print(grid_search_results)

# Elegimos el mejor modelo
best_params = {}
for m in grid_search_results:
    if m['classification_error'] < best_params.get('classification_error', np.inf):
        best_params = m
print(best_params)

0%                                100%
[#####] | ETA: 00:00:00

[{'activation_func': 'tanh', 'epochs': 10, 'optimizer_method': 'SGD', 'l1_reg': 1, 'hidden': [64, 32], '
{'activation_func': 'tanh', 'epochs': 10, 'optimizer_method': 'adadelata', 'l1_reg': 0.0001, 'hidden': [6

Total time elapsed: 00:43:51

In [13]: d = [{'activation_func': 'tanh', 'epochs': 10, 'optimizer_method': 'SGD', 'hidden': [64, 32],

```

3.0.3 Entrenamiento del modelo final

```

In [14]: model = create_model(best_params)
         training_results = model.fit(np_X_train, np_Y_train,
                                     validation_data=(np_X_test, np_Y_test),
                                     nb_epoch=10)

```

Train on 7352 samples, validate on 2947 samples

Epoch 1/10

7352/7352 [=====] - 3s - loss: 0.8128 - acc: 0.8104 - val_loss: 0.3914 - val_acc: 0.8104

Epoch 2/10

7352/7352 [=====] - 2s - loss: 0.4567 - acc: 0.9144 - val_loss: 0.2308 - val_acc: 0.9144

Epoch 3/10

7352/7352 [=====] - 2s - loss: 0.3675 - acc: 0.9433 - val_loss: 0.2580 - val_acc: 0.9433

Epoch 4/10

7352/7352 [=====] - 2s - loss: 0.3203 - acc: 0.9546 - val_loss: 0.1705 - val_acc: 0.9546

Epoch 5/10

7352/7352 [=====] - 2s - loss: 0.2970 - acc: 0.9601 - val_loss: 0.2073 - val_acc: 0.9601

Epoch 6/10

7352/7352 [=====] - 2s - loss: 0.2773 - acc: 0.9642 - val_loss: 0.2663 - val_acc: 0.9642

Epoch 7/10

7352/7352 [=====] - 2s - loss: 0.2612 - acc: 0.9701 - val_loss: 0.1400 - val_acc: 0.9701

Epoch 8/10

7352/7352 [=====] - 2s - loss: 0.2495 - acc: 0.9718 - val_loss: 0.1389 - val_acc: 0.9718

Epoch 9/10


7352/7352 [=====] - 2s - loss: 0.2358 - acc: 0.9744 - val_loss: 0.1907 - val_acc: 0.9744

Epoch 10/10

7352/7352 [=====] - 2s - loss: 0.2314 - acc: 0.9732 - val_loss: 0.1569 - val_acc: 0.9732

```
In [15]: pylab.plot(training_results.epoch, training_results.history['acc'], label='training')
pylab.plot(training_results.epoch, training_results.history['val_acc'], label='test')
pylab.legend(loc='lower right')
pylab.xlabel(u"Época")
pylab.ylabel(u"Precisión")
pylab.title(u"Precisión de clasificación vs época")
```

Out[15]: <matplotlib.text.Text at 0x7fe455a3b5d0>



4 Evaluating models

```
In [37]: error = 1 - model.evaluate(np_X_train, np_Y_train)[1]
         print("\nE_in = {}".format(error))
```

```
7328/7352 [=====>.] - ETA: 0s
E_in = 0.0214907508161
```

```
In [38]: error = 1 - model.evaluate(np_X_test, np_Y_test)[1]
         print("\nE_test = {}".format(error))
```

```
2912/2947 [=====>.] - ETA: 0s
E_test = 0.061757719715
```

```
In [41]: from sklearn.metrics import confusion_matrix
         classes = model.predict_classes(np_X_test)
         print "\nMatriz de confusión:"
         print confusion_matrix(np.argmax(np_Y_test,axis=1), classes).tolist()
```

```
2912/2947 [=====>.] - ETA: 0s
```

Matriz de confusión:

```
[[526, 0, 11, 0, 0, 0], [0, 393, 96, 0, 0, 2], [0, 14, 518, 0, 0, 0], [0, 0, 0, 482, 9, 5], [0, 0, 0, 4
```